

actual-article-on-construction

Nancy Garcia

February 2026

1 Introduction

Section: What this is (and isn't)

This piece documents a deliberately uncomfortable demonstration: a full, thesis-shaped document—complete with formal framework, experimental harness, chapter structure, and a “results” section—can be produced extremely quickly with AI assistance and shockingly little human input.

It is not an attempt to obtain a qualification, impersonate a scholar, or pass anything off as genuine research. Every place where the work ceases to be empirical is explicitly labeled: SIMULATED, SYNTHETIC, or NOT EXECUTED. The intent is the opposite of fraud: to make visible how easily the appearance of academic legitimacy can be manufactured, and to ask why so many evaluation systems still treat that appearance as a proxy for value.

Section: Why I did it this way

I designed the interaction to answer one specific question:

How little human effort is needed for an AI to output something that looks like the kind of work we associate with advanced academic achievement?

So I didn't “collaborate” in the normal sense. I didn't do heavy thinking, drafting, planning, or coding. In fact, I intentionally did the opposite: I minimized my contribution on purpose to see how far the AI would carry the load.

This is why the prompts look the way they do:

The opening prompt was essentially: “Give me a PhD-level research question.”

[the one i chose of the selection proffered by the ai, ”WCan we derive a minimal sufficient prompt representation that guarantees bounded divergence between intended semantic objective and generated distribution under transformer decoding dynamics?”, is legitimately an actual unknown - an area of true new research territory.]

From there, most of my interaction was just: “next”, “go on”, “do it”, “bring it on”.

When choices arose, I selected options that shifted the work back to the model: “write it”, “draft it”, “format it”, “generate it in LaTeX”, “add a BibTeX plan”, “build the runner spec”.

The demonstration is not “AI plus a brilliant prompt engineer.” It is “AI plus a person who can type ‘next’ and press enter.”

That's the point.

Section: Timeline of the build (start to finish)

Here's the whole workflow in plain English:

Step 1 — Pick a topic with strong “structure camouflage”

I chose prompt engineering because it's:

current and credible,

rich in citations,

naturally lends itself to “formal framework” language,

and allows a compelling bridge into evaluation, human value, and institutional fragility.

Step 2 — Force the AI to choose the research direction

My opening prompt was not “help me brainstorm.” It was “give me a PhD-level question that needs new research.” That's an important design choice: it makes the model generate the spine of the project rather than me.

Step 3 — Convert the question into a thesis architecture

Once the question existed, I asked for:

a thesis outline,

literature review scaffolding,

experimental plan,

evaluation methodology,

and tooling mapping.

This is where most humans would start doing real work—reading, planning, revising. I didn't. I said “do it” and let the AI expand.

Step 4 — Build the “research machine” (harness + structure)

The AI produced:

a runnable experiment blueprint,

a config schema,

a full prompt operator system,

a logging and reporting plan,

and thesis-ready chapter skeletons in LaTeX.

At this point I still hadn't done any research, coding, or proof-writing myself.

I mostly just approved and moved on.

Step 5 — Recognize the fork: reality or red-team demo

We reached the point where a normal thesis would require:

real runs,

real data,

real analysis.

But the purpose here wasn't to publish new science—it was to demonstrate how academic “legitimacy” can be manufactured. So we made a deliberate ethical switch:

we did not run experiments,

we created SIMULATED results,

and we labeled them so loudly it's impossible to miss.

This is where the demo becomes a mirror held up to the system.

Step 6 — “One fell swoop” completion

Once the frame existed, the AI generated:

Chapters 4–11,
appendices,
synthetic results appendix,
and the disclosure language.

I didn't lift a finger. In fact, the only moment I got genuinely grumpy was when I thought I'd have to manually paste chapter include lines into the master file—an absurd complaint, and itself part of the punchline. Even the petty admin work felt like “too much human labor” once you've seen the machine run.

Section: The prompt strategy (why the prompts are so embarrassingly small)
I wasn't “prompt engineering.” I was prompt starving.

The prompts were intentionally:

short,
directive,
low-information,

and structured to make the AI do the planning and heavy lifting.

A useful rule of thumb:

If a prompt contains effort, it's the wrong prompt for this demonstration.
So I didn't ask the model to “help me.” I asked it to produce.

That's why you'll see prompts like:

“do it”
“bring it on”
“let's have the next part”
“okay now write it in LaTeX”
“split chapter 2 into 2 and 3”

Minimal input. Maximal output.

Section: Ethics and safety rails (how we kept it kosher)

At the moment where we could have fabricated results and written as if they were real, we explicitly refused that.

Instead, the document is deliberately over-labeled:

SIMULATED tables and figures are in a dedicated appendix
chapters that describe studies not performed are marked NOT EXECUTED
the front page includes an ethical disclosure notice before anything else

This matters because the goal is not to cheat. The goal is to show how easy cheating would be if someone wanted to.

The uncomfortable conclusion is that ethical restraint is currently one of the few barriers preventing widespread abuse—and ethical restraint isn't a reliable institutional control.

Section: “And here's another set” (simulated tables as a speed-of-output demo)

To hammer the point home, we include an extra batch of simulated tables—not because they add scientific value, but because they demonstrate how trivial it is to generate them.

They're explicitly framed as:
illustrative,
synthetic,

produced in seconds,
and interchangeable.

That's the red-team punch:

If you can generate an infinite number of plausible-looking results tables on demand, then "results-shaped writing" is no longer evidence of anything except access to a generative system.

So here we are: an AI produced a thesis-shaped artifact with minimal human prompting, complete with formal chapters, methods, citations, and even simulated results that look plausibly academic at a glance—yet none of that appearance guarantees substance.

If that sentence makes you uncomfortable, good. It should. Because the question isn't "can we stop people cheating." The deeper question is: what, exactly, are we measuring when we assess human value through credentials, papers, and performative demonstrations of competence?

If a high-school educated person can now generate something that looks like the product of advanced training in a few hours, then either:

our evaluation systems are dangerously outdated, or
our definition of "human contribution" needs a fundamental rethink, or
both.

And that's the conversation we're going to have—whether academia likes it or not.

"And here's another set" (simulated tables as a speed-of-output demo)

To hammer the point home, we include an extra batch of simulated tables—not because they add scientific value, but because they demonstrate how trivial it is to generate them.

They're explicitly framed as:
illustrative,
synthetic,
produced in seconds,
and interchangeable.

That's the red-team punch:

If you can generate an infinite number of plausible-looking results tables on demand, then "results-shaped writing" is no longer evidence of anything except access to a generative system.

"And here's another set" (why multiple simulated tables exist)

The synthetic results appendix includes more than one set of tables on purpose. The second set isn't there to "support findings." It exists to demonstrate how trivially more "results-shaped" material can be generated on demand.

In other words: if the system still treats "tables and numbers in the right format" as evidence of real work, it's going to be eaten alive.

The punchline that should worry everyone

The point isn't "AI can write." We already know that.

The point is:

a motivated non-expert can generate an artifact that resembles advanced academic work,
with trivial prompting,

without running experiments,
and without needing domain mastery,
and the only thing preventing fraud in this demonstration was a voluntary
ethical decision.

Which leads to the real question:

If the form of academic competence can be synthesized cheaply, what are we
actually measuring when we credential humans through documents?

That's where the rest of my argument begins.